

Session 3-Assignment 3 -YARN

Task 1:

*Execute WordMedian , WordMean , WordStandardDeviation programs using
hadoop-mapreduce-examples-2.9.0.jar file present in your Acadgild VM*

Solution :

1)Executing wordmedian program

*First create a input text file “text.txt” with some data in it and place the file in
/user/hadoopdata/hdfs.*

*Now change the directory to mapreduce where the jar hadoop-mapreduce-examples-2.6.5.jar is
stored using the below command:*

cd/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce/

Executed wordmedian program using below command.

***hadoop jar hadoop-mapreduce-examples-2.6.5.jar wordmedian
/user/hadoopdata/hdfs/text.txt /user/hadoopdata/hdfs/wordmedian_output***

Below is the screenshot of the wordmedian program output

```
Applications Places System Tue May 15, 11:00 PM Acadgild
acadmild@localhost:~/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce
File Edit View Search Terminal Help
[acadmild@localhost ~]$ cd /home/acadmild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce/
[acadmild@localhost mapreduce]$ hadoop jar hadoop-mapreduce-examples-2.6.5.jar wordmedian /user/hadoopdata/hdfs/text.txt /user/hadoopdata/hdfs/Wordmedian_output
18/05/15 22:58:03 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/05/15 22:58:04 INFO client.RMPProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/05/15 22:58:06 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
18/05/15 22:58:06 INFO input.FileInputFormat: Total input paths to process : 1
18/05/15 22:58:06 INFO mapreduce.JobSubmitter: number of splits:1
18/05/15 22:58:07 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1526401988903_0002
18/05/15 22:58:07 INFO impl.YarnClientImpl: Submitted application application_1526401988903_0002
18/05/15 22:58:07 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1526401988903_0002/
18/05/15 22:58:07 INFO mapreduce.Job: Running job: job_1526401988903_0002
18/05/15 22:58:23 INFO mapreduce.Job: Job job_1526401988903_0002 running in uber mode : false
18/05/15 22:58:23 INFO mapreduce.Job: map 0% reduce 0%
18/05/15 22:58:36 INFO mapreduce.Job: map 100% reduce 0%
18/05/15 22:58:59 INFO mapreduce.Job: map 100% reduce 100%
18/05/15 22:59:00 INFO mapreduce.Job: Job job_1526401988903_0002 completed successfully
18/05/15 22:59:00 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=66
    FILE: Number of bytes written=215275
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=160
    HDFS: Number of bytes written=24
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=11501
    Total time spent by all reduces in occupied slots (ms)=19839
    Total time spent by all map tasks (ms)=11501
    Total time spent by all reduce tasks (ms)=19839
    Total vcore-milliseconds taken by all map tasks=11501
    Total vcore-milliseconds taken by all reduce tasks=19839
    Total megabyte-milliseconds taken by all map tasks=11777024
    Total megabyte-milliseconds taken by all reduce tasks=20315136
  Map-Reduce Framework
    Map input records=4
    Map output records=7
    Map output bytes=56
    Map output materialized bytes=66
    Input split bytes=116
    Combine input records=7
    Combine output records=6
    Reduce input groups=6
    Reduce shuffle bytes=66
    Reduce input records=6
    Reduce output records=6
    Spilled Records=12
    Shuffled Maps =1
    Failed Shuffles=0
    Merged Map outputs=1
    GC time elapsed (ms)=233
    CPU time spent (ms)=2470
    Physical memory (bytes) snapshot=295452672
    Virtual memory (bytes) snapshot=4118224896
    Total committed heap usage (bytes)=170004480
  Shuffle Errors
    BAD_ID=0
    CONNECTION=0
```

```

IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
    Bytes Read=44
File Output Format Counters
    Bytes Written=24
The median is: 6
You have new mail in /var/spool/mail/acadgild

```

2)Executing wordmean program

Change the directory to mapreduce where the jar `hadoop-mapreduce-examples-2.6.5.jar` is stored using the below command:

```
cd /home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce/
```

Execute wordmean program using below command.

```
hadoop jar hadoop-mapreduce-examples-2.6.5.jar wordmean  
/user/hadoopdata/hdfs/text.txt /user/hadoopdata/hdfs/wordmean_output
```

Below is the screenshot of the wordmean program output

```

[acadgild@localhost ~]$ cd /home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce/
[acadgild@localhost mapreduce]$ hadoop jar hadoop-mapreduce-examples-2.6.5.jar wordmean /user/hadoopdata/hdfs/text.txt /user/hadoopdata/hdfs/Wordmean.output
18/05/15 23:09:55 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/05/15 23:09:56 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/05/15 23:09:58 INFO input.FileInputFormat: Total input paths to process : 1
18/05/15 23:09:58 INFO mapreduce.JobSubmitter: number of splits:1
18/05/15 23:09:58 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1526401988903_0003
18/05/15 23:09:59 INFO impl.YarnClientImpl: Submitted application application_1526401988903_0003
18/05/15 23:09:59 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1526401988903_0003/
18/05/15 23:09:59 INFO mapreduce.Job: Running job: job_1526401988903_0003
18/05/15 23:10:13 INFO mapreduce.Job: Job job_1526401988903_0003 running in uber mode : false
18/05/15 23:10:13 INFO mapreduce.Job:  map 0% reduce 0%
18/05/15 23:10:23 INFO mapreduce.Job:  map 100% reduce 0%
18/05/15 23:10:34 INFO mapreduce.Job:  map 100% reduce 100%
18/05/15 23:10:34 INFO mapreduce.Job: Job job_1526401988903_0003 completed successfully
18/05/15 23:10:34 INFO mapreduce.Job: Counters: 49
    File System Counters
      FILE: Number of bytes read=39
      FILE: Number of bytes written=215481
      FILE: Number of read operations=0
      FILE: Number of large read operations=0
      FILE: Number of write operations=0
      HDFS: Number of bytes read=160
      HDFS: Number of bytes written=18
      HDFS: Number of read operations=6
      HDFS: Number of large read operations=0
      HDFS: Number of write operations=2
    Job Counters
      Launched map tasks=1
      Launched reduce tasks=1
      Data-local map tasks=1
      Total time spent by all maps in occupied slots (ms)=8032

```

```

Total time spent by all reduces in occupied slots (ms)=7810
Total time spent by all map tasks (ms)=8032
Total time spent by all reduce tasks (ms)=7810
Total vcore-milliseconds taken by all map tasks=8032
Total vcore-milliseconds taken by all reduce tasks=7810
Total megabyte-milliseconds taken by all map tasks=8224768
Total megabyte-milliseconds taken by all reduce tasks=7997440
Map-Reduce Framework
  Map input records=4
  Map output records=14
  Map output bytes=203
  Map output materialized bytes=39
  Input split bytes=116
  Combine input records=14
  Combine output records=2
  Reduce input groups=2
  Reduce shuffle bytes=39
  Reduce input records=2
  Reduce output records=2
  Spilled Records=4
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=277
  CPU time spent (ms)=2160
  Physical memory (bytes) snapshot=297414656
  Virtual memory (bytes) snapshot=4118228992
  Total committed heap usage (bytes)=170004480
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=44
File Output Format Counters
  Bytes Written=18
The mean is: 5.285714285714286
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost mapreduce]$

```

3)Executing WordStandardDeviation program

Change the directory to mapreduce where the jar *hadoop-mapreduce-examples-2.6.5.jar* is stored using the below command:

cd/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce/

Execute wordstandardeviation mean program using below command.

***hadoop jar hadoop-mapreduce-examples-2.6.5.jar wordstandarddeviation
/user/hadoopdata/hdfs/text.txt /user/hadoopdata/hdfs/WSDevaition_output***

Below is the screenshot of the wordstandrddeviation program output

File Edit View Search Terminal Help

```
[acadgild@localhost ~]$ cd /home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce/
[acadgild@localhost mapreduce]$ hadoop jar hadoop-mapreduce-examples-2.6.5.jar wordstandarddeviation /user/hadoopdata/hdfs/text.txt /user/hadoopdata/hdfs/WSDeviation.output
18/05/15 23:19:58 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/05/15 23:20:00 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/05/15 23:20:02 INFO input.FileInputFormat: Total input paths to process : 1
18/05/15 23:20:02 INFO mapreduce.JobSubmitter: number of splits:1
18/05/15 23:20:02 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1526401988903_0004
18/05/15 23:20:03 INFO impl.YarnClientImpl: Submitted application application_1526401988903_0004
18/05/15 23:20:03 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1526401988903_0004/
18/05/15 23:20:03 INFO mapreduce.Job: Running job: job_1526401988903_0004
18/05/15 23:20:16 INFO mapreduce.Job: Job job_1526401988903_0004 running in uber mode : false
18/05/15 23:20:16 INFO mapreduce.Job: map 0% reduce 0%
18/05/15 23:20:26 INFO mapreduce.Job: map 100% reduce 0%
18/05/15 23:20:36 INFO mapreduce.Job: map 100% reduce 100%
18/05/15 23:20:37 INFO mapreduce.Job: Job job_1526401988903_0004 completed successfully
18/05/15 23:20:37 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=56
    FILE: Number of bytes written=215681
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=160
    HDFS: Number of bytes written=29
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=6940
    Total time spent by all reduces in occupied slots (ms)=7712
    Total time spent by all map_tasks (ms)=6940

    Total time spent by all reduces in occupied slots (ms)=7712
    Total time spent by all map tasks (ms)=6940
    Total time spent by all reduce tasks (ms)=7712
    Total vcore-milliseconds taken by all map tasks=6940
    Total vcore-milliseconds taken by all reduce tasks=7712
    Total megabyte-milliseconds taken by all map tasks=7106560
    Total megabyte-milliseconds taken by all reduce tasks=7897088
  Map-Reduce Framework
    Map input records=4
    Map output records=21
    Map output bytes=308
    Map output materialized bytes=56
    Input split bytes=116
    Combine input records=21
    Combine output records=3
    Reduce input groups=3
    Reduce shuffle bytes=56
    Reduce input records=3
    Reduce output records=3
    Spilled Records=6
    Shuffled Maps =1
    Failed Shuffles=0
    Merged Map outputs=1
    GC time elapsed (ms)=212
    CPU time spent (ms)=2370
    Physical memory (bytes) snapshot=296300544
    Virtual memory (bytes) snapshot=4123373568
    Total committed heap usage (bytes)=170004480
  Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
  File Input Format Counters
    Bytes Read=44
```

File Output Format Counters

Bytes Written=29

The standard deviation is: 1.9794866372215743

You have new mail in /var/spool/mail/acadgild

[acadgild@localhost mapreduce]\$